

The code below is a companion to the paper:

J. Chakravorty, J. Subramanian, A. Mahajan, “Stochastic approximation based methods for computing the optimal thresholds in remote-state estimation with packet drops,” in proceedings of the American Control Conference, Seattle, WA, 2017.

## Introduction

The code below performs the following calculations. Consider a first-order auto-regressive process  $\{E_t\}_{t \geq 0}$  which evolves as follows:

$$E_{t+1} = aE_t + W_t,$$

where  $a$  is a real number and  $\{W_t\}_{t \geq 0}$  is a i.i.d. process with symmetric and unimodal distribution. In our simulations, we restrict attention to  $a = 1$  and  $W_t$  distributed according to zero-mean Gaussian distribution with unit variance. These values are hard-coded in the `nextState` function defined below.

```
@inline function nextState(E)
    a = 1.0
    σ = 1.0
    a*E + σ*randn()
end
```

Whenever  $|E_t| < k$ , where  $k$  is a threshold that can be tuned, we incur a distortion  $d(E_t)$ . We hard-code  $d(e) = e^2$  in the distortion function defined below.

```
@inline distortion(E) = E^2
```

Let  $S_t$  denote the state of the channel. It is assumed that  $\{S_t\}_{t \geq 0}$  is an i.i.d. Bernoulli process with  $\Pr(S_t = 0) = p_d$ .

## Sampling functions

Let  $\tau$  denote the stopping time when  $\{|E_t| \geq k\}$  and  $\{S_t = 1\}$ . The code below performs the following sample-path calculations.

$$L = \sum_{t=0}^{\tau-1} \beta^t d(E_t), \quad M = \sum_{t=0}^{\tau-1} \beta^t, \quad K = \sum_{t=0}^{\tau} \beta^t U_t,$$

where  $\beta$  is the discount factor and  $U_t = \mathbb{I}\{|E_t| \geq k\}$ .

Since the calculation is stochastic, there is a positive probability that  $\tau$  is a big number, which can slow down the calculations. So, we set a bound `maxIterations` on the maximum size of  $\tau$ . We set the

default value of `maxIterations` as `10_000`. This number may need to be increased when computing the performance for large threshold.

```
@inline function sample(threshold, discount, dropProb; maxIterations = 10_000)

    E, L, M, K = 0.0, 0.0, 0.0, 0.0

    count = 0
    scale = 1.0

    while count <= maxIterations
        channel_on = rand() > dropProb
        transmit    = !(-threshold < E < threshold)
        success     = transmit && channel_on

        if !success
            L += scale * distortion(E)
            K += scale * transmit
            M += scale
        else
            K += scale * transmit
            # Note that we could have written
            #     k += scale
            # because `transmit` is true in this branch of the code.
            break
        end

        scale *= discount
        E = nextState(E)
        count += 1
    end

    (L, M, K)
end
```

### Mini-batch averaging

The value obtained by one sample path is usually noisy. So, we smoothen it out by averaging over a mini-batch. The default size of the mini-batch is 1000 iterations.

```

@inline function sample_average(threshold, discount, dropProb; iterations::Int=1000)

    ell, emm, kay = 0.0, 0.0, 0.0

    for i in 1:iterations
        L, M, K = sample(threshold, discount, dropProb)
        ell += L
        emm += M
        kay += K
    end
    ell /= iterations
    emm /= iterations
    kay /= iterations
    (ell, emm, kay)
end

```

## Stochastic approximation

### Stochastic gradient descent for costly communication

It has been shown in the paper that a threshold  $k$  is optimal for communication cost  $\lambda$  if

$$\partial_k C(k) = 0, \quad \text{where } C(k) = D(k) + \lambda N(k), \quad D(k) = \frac{L(k)}{M(k)}, \quad N(k) = \frac{K(k)}{M(k)}.$$

The function `sa_costly` computes the optimal threshold for given values of `cost`, `discount`, and `dropProb`. It uses Kiefer Wolfowitz algorithm. In particular, the gradient is calculated using finite differences:

$$\nabla L \approx \frac{1}{2c} [L(k+c) - L(k-c)].$$

By default, `c` is set to `0.1`. If we were writing this code for higher dimensions, we would replace Kiefer-Wolfowitz with the simultaneous perturbation (SPSA) algorithm, which is more sample efficient for higher dimensions.

The stochastic approximation iteration starts from an initial guess (the parameter `initial`). Its default value is `1.0`. This initialization could be useful if we have a reasonable guess for optimal threshold (e.g., the exact solution obtained by Fredholm integral equations for the case when `dropProb` = `0`; see TAC 2017 paper for details).

It is not possible to detect convergence of stochastic approximation algorithms. So we run the algorithm for a fixed number of iterations (the parameter `iterations`, whose default value is `1_000`).

Stochastic approximation algorithms are sensitive to the choice of learning rates. We use ADAM to adapt the learning rates according to the sample path. The parameters `decay1`, `decay2`, `alpha`, and `epsilon` can be used to tune ADAM. In our experience, these should be left to their default values.

It is not possible to detect convergence of stochastic approximation algorithms. So we run the algorithm for a fixed number of iterations (the parameter `iterations`, whose default value is `1_000`).

Sometimes it is useful to visualize the estimates (of the threshold) as the algorithm is running. To do so, set `debug` to `true`, which will print the current estimate of the threshold after every 100 iterations.

The output of the function is a trace of the estimates of the threshold (therefore, it is a 1D array of size `iterations`).

```
@fastmath function sa_costly(cost, discount, dropProb ;
    iterations :: Int      = 1_000,
    initial    :: Float64 = 1.0,
    decay1     :: Float64 = 0.9,
    decay2     :: Float64 = 0.999,
    epsilon    :: Float64 = 1e-8,
    alpha      :: Float64 = 0.01,
    c          :: Float64 = 0.1,
    debug      :: Bool     = false,
)

    threshold = initial
    trace      = zeros(iterations)

    moment1 = 0.0
    moment2 = 0.0

    weight1 = decay1
    weight2 = decay2

    @inbounds for k in 1:iterations
        threshold_plus = threshold + c
        threshold_minus = threshold - c

        L_plus , M_plus, K_plus = sample_average(threshold_plus, discount, dropProb)
        L_minus, M_minus, K_minus = sample_average(threshold_minus, discount, dropProb)

        C_plus  = (L_plus  + cost*K_plus )/M_plus
        C_minus = (L_minus + cost*K_minus)/M_minus
```

```

gradient = (C_plus - C_minus)/2c

moment1 = decay1 * moment1 + (1 - decay1) * gradient
moment2 = decay2 * moment2 + (1 - decay2) * gradient^2

corrected1 = moment1/(1 - weight1)
corrected2 = moment2/(1 - weight2)

weight1 *= decay1
weight2 *= decay2

threshold_delta = corrected1 / ( sqrt(corrected2) + epsilon)

threshold -= alpha * threshold_delta

# The above analysis is valid for positive values of threshold, but
# the above line can make threshold negative. Ideally, if the
# threshold is negative, we should set it to zero. But if the
# threshold is set to zero, the value of M in the next iteration will
# also be zero; therefore, the C at the next instant will be infinity.
# So, we set the minimum value of threshold to be 2c.

threshold = max(threshold, 2c)

if debug && mod(k,100) == 0
    @printf("#:%8d, threshold=%0.6f\n", k, threshold)
end

trace[k] = threshold
end

return trace
end

```

### Stochastic gradient descent for constrained communication

It has been shown in the paper that a threshold  $k$  is optimal for a communication rate  $\alpha$  if

$$\alpha M(k) - K(k) = 0.$$

The function `sa_constrained` computes the optimal threshold for a given value of `rate`, `discount`, and `dropProb`. It uses Robbins-Monro algorithm. This algorithm converges fairly fast, so we do not use mini-batches to smoothen out the samples and set the learning rate at iteration  $n$  to be  $\alpha/n$ , where  $\alpha$  is set to 1, by default.

Sometimes it is useful to visualize the estimates (of the threshold) as the algorithm is running. To do so, set `debug` to `true`, which will print the current estimate of the threshold after every 100 iterations.

The output of the function is a trace of the estimates of the threshold (therefore, it is a 1D array of size `iterations`).

```
@fastmath function sa_constrained(rate, discount, dropProb ;
    iterations :: Int      = 1_000,
    initial :: Float64 = 1.0,
    alpha  :: Float64 = 1.0,
    debug  :: Bool     = false,
)

    threshold = initial
    trace      = zeros(iterations)

    @inbounds for k in 1:iterations
        _, M, K = sample_average(threshold, discount, dropProb; iterations=1)

        rl      = alpha/k
        threshold -= rl*(rate * M - K)

        # As in the case of costly communication, the above line can make the
        # threshold negative, so we set its minimum value to be a small
        # positive number (0.05). Note that hard-coding this value means that
        # the code will not work correctly for really small values of the
        # communication rate, where we expect the optimal threshold to be
        # less than 0.05.

        threshold = max(threshold, 0.05)

        if debug && mod(k,100) == 0
            @printf("#:%8d, threshold=%0.6f\n", k, threshold)
        end

        trace[k] = threshold
    end
end
```

```
end
```

```
return trace
```

```
end
```